

# AI Jukebox

---

*Testing Creativity in AI*

Brian McMahon

5 April 2018

# ***Business Case***



**Exploring the question: can AI be creative?**


- Expands scope of artistic applications in the arts, both audio and visual

**Audio analysis and implementation essential for:**

- Digital assistants
- Speech to text

**Application of deep learning:** testing the boundaries of AI ability to positively impact the human experience

# *Tools*



music21



# Dataset



## Midi files by genre

- Final Fantasy soundtrack:
  - 91 midi files, 51,177 notes and 358 unique notes
  - Successfully utilized in previous successful music generator models\*
- Celtic folk tunes:
  - 338 midi files, 159,789 notes and 78 unique notes
  - Distinct, upbeat

\*See [blog post](#) by Sigurour Skuli, Towards Data Science.

# LSTM Network



- Have “memory”, allowing information to persist, including *long-term dependencies*
- At each timestep, previous state is passed in along with new input
- “Gate” functionality: input, cell/forget, output


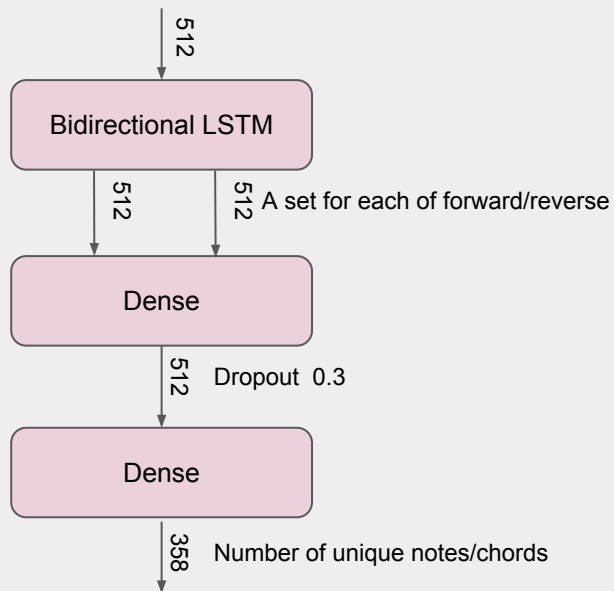


Diagram courtesy of ["colah's blog"](#).

# Architecture



## Bidirectional LSTM



- 512 node input layer
- Dropout on dense layer only
- Learning rate 0.001
- Sequence length 200
- Notes generated 500

# Generated Notes

*Output from the artificial neural network*



Robot image courtesy of [trendhunter.com](https://www.trendhunter.com).

# Results



**As the model is generative, the best judges are us!**

Looking for:

- Sensible recurring patterns, melodic
- Pleasing to the ear
- Model generates low loss on data, and validation loss on test set

**To put into practice:**

- Utilize deep learning and neural networks in interpreting audio data
- Generate new content which, perhaps, is removed from human environment
- Supplement artistic work with new rhythms, beats and patterns generated by AI



# ***Next Steps***



**Continue to refine model performance.** Explore:

- Different datasets
- different architectures, such as variational autoencoders
- Different inputs, such as raw audio

**Write model into flask app and implement online**

- Input a set of midi files, output AI music!

# Thank You!

*bcm822@gmail.com*

# Appendix

# ***Lessons Learned***



- Successfully implemented a functional AI music generator
- Tested the audio and generative capabilities of neural networks
- Utilized various audio format preprocessing



# Resources

Dorsey, Brannon. "Using Machine Learning to Create New Melodies.." <https://brangerbriz.com/>. 10 May 2017.

Nayebi, Aran. "GRUV: Algorithmic Music Generation using Recurrent Neural Networks." Stanford University. 2015.

Skúli, Sigurður. "How to Generate Music using a LSTM Neural Network in Keras." [www.towardsdatascience.com](http://www.towardsdatascience.com). December 7, 2017.

Brownlee, Jason. "Stacked LSTM Networks." <https://machinelearningmastery.com>. August 18, 2017.

Brownlee, Jason. "Understand the Difference Between Return Sequences and Return States for LSTMs in Keras." <https://machinelearningmastery.com>. October 24, 2017.

"Understanding LSTM Networks." Colah's Blog. <https://colah.github.io>. 27 August 2015.





# Sequence Generation

A B C D E



Network



B C D E F



Network




C D E F G

- Model generates each note/chord by looking at the previous 100 and taking the highest probability next note/chord
- This shifts the considered set by 1 each time

Classes

A B C D E F G




Network Output

0.6 0.34 0.1 0.92 0.47 0.22 0.69



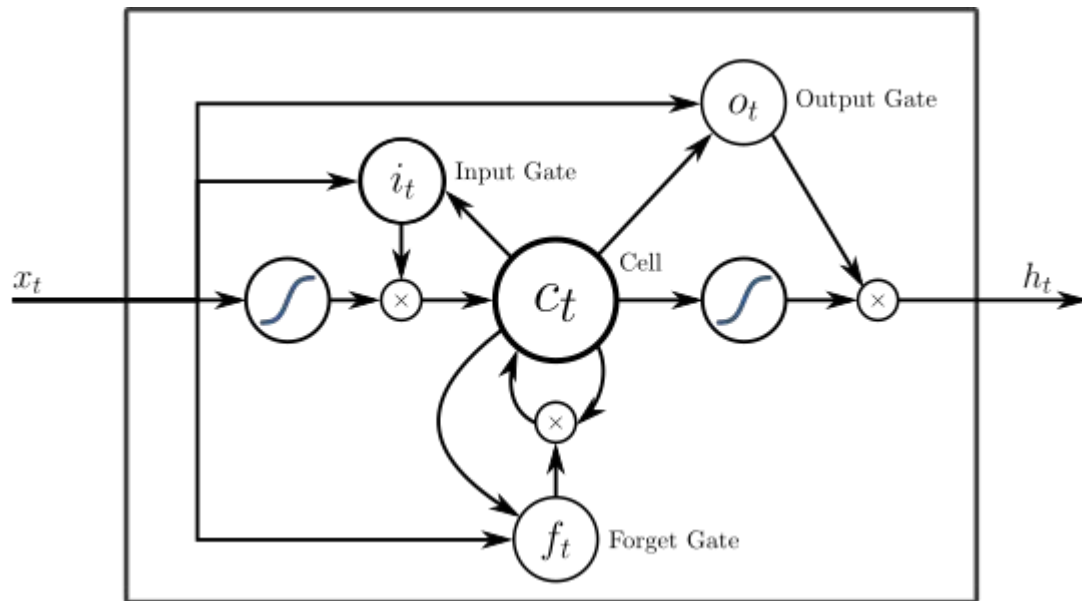


# ***Recurrent Network***





# ***LSTM Diagram (2)***







# Model (2)

In [5]: `model.summary()`

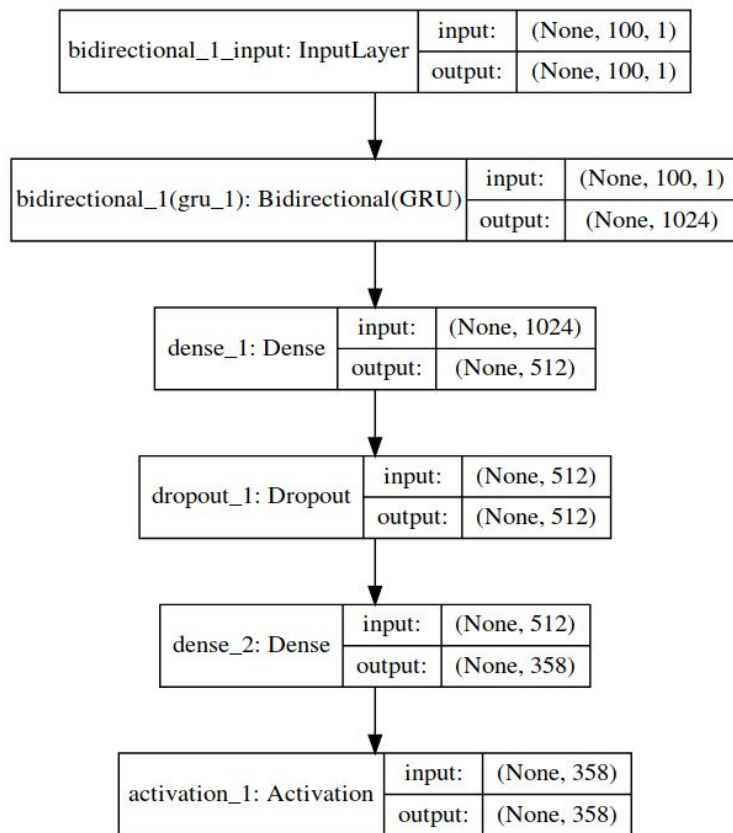
Layer (type)	Output Shape	Param #
=====		
bidirectional_1 (Bidirection	(None, 1024)	1579008
dense_1 (Dense)	(None, 512)	524800
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 358)	183654
activation_1 (Activation)	(None, 358)	0
=====		
Total params: 2,287,462		
Trainable params: 2,287,462		
Non-trainable params: 0		





# Model (3)

[note this is GRU]





# Model (4)

```
model = Sequential()
model.add(Bidirectional(LSTM(first_layer), input_shape=(timesteps, data_dim)))
model.add(Dense(first_layer))
model.add(Dropout(drop))
model.add(Dense(n_vocab)) # based on number of unique notes
model.add(Activation('softmax'))

rms = optimizers.RMSprop(lr=0.001, rho=0.9, epsilon=None, decay = 0.0)
model.compile(loss='categorical_crossentropy', optimizer=rms)
```



